


- 
- Probability Primer
 - Naïve Bayes
 - ▣ Bayes' Rule
 - ▣ Conditional Probabilities
 - ▣ Probabilistic Models

Motivation

- In many datasets, relationship between attributes and a class variable is non-deterministic.
- Why?
 - ▣ Noisy data
 - ▣ Confounding and interaction of factors
 - ▣ Relevant variables not included in the data
- *Scenario*:
 - ▣ Risk of heart disease based on individual's diet and workout frequency

Scenario

- Risk of heart disease based on individual's diet and workout frequency
 - ▣ Most people who “work out” and have a healthy diet don't get heart disease
 - ▣ Yet, some healthy individuals still do:
 - Smoking, alcohol abuse, ...

What we're trying to do

- Model probabilistic relationships
- “What is the *probability that this person will get heart disease, given their diet and workout regimen?*”
 - ▣ Output is most similar to Logistic Regression
- Will introduce naïve Bayes model
 - ▣ A type of Bayesian classifier
 - ▣ More advanced: Bayesian network

Bayes Classifier

- A probabilistic framework for solving *classification* problems
 - ▣ Used in both naïve Bayes and Bayesian networks
- Based on Bayes' Theorem:

What's the probability that it rains today *AND* that I'm carrying an umbrella?

Terminology/Notation Primer

- X and Y (two different variables)
- Joint probability: $P(X=x, Y=y)$
 - The probability that variable X takes on the value x and variable Y has the value y
- Conditional probability: $P(Y=y \mid X=x)$
 - Probability that variable Y has the value y , given that variable X takes on the value x

Given that I'm observed with an umbrella, what's the probability that it will rain today?

Terminology/Notation Primer

- Single Probability: $P(X = x)$
 - ▣ “X has the value x”
- Joint Probability: $P(X, Y)$
 - ▣ “X and Y”
- Conditional Probability: $P(Y | X)$
 - ▣ “Y” given observation of “X”
- Relation of Joint and Conditional Probabilities:

$$P(X, Y) = P(Y | X) \cdot P(X)$$

Terminology/Notation Primer

$$P(X, Y) = P(Y | X) \cdot P(X)$$

$$P(Y, X) = P(X | Y) \cdot P(Y)$$

$$P(X, Y) = P(Y, X)$$

$$P(X, Y) = P(Y | X) \cdot P(X) = P(X | Y) \cdot P(Y)$$

Bayes' Theorem:
$$P(Y | X) = \frac{P(X | Y) \cdot P(Y)}{P(X)}$$

Predicted Probability Example

□ Scenario:

1. A doctor knows that meningitis causes a stiff neck 50% of the time
 2. Prior probability of any patient having meningitis is $1/50,000$
 3. Prior probability of any patient having a stiff neck is $1/20$
- If a patient has a stiff neck, what's the probability that they have meningitis?

- A doctor knows that meningitis causes a stiff neck 50% of the time
- Prior probability of any patient having meningitis is $1/50,000$
- Prior probability of any patient having a stiff neck is $1/20$

□ Apply Bayes' Rule:

□ If a patient has a stiff neck, what's the probability that they have meningitis?

■ Interested in: $P(M | S)$

$$P(M | S) = \frac{P(S | M)P(M)}{P(S)}$$

$$P(M | S) = \frac{P(S | M)P(M)}{P(S)} = \frac{0.5 \cdot \frac{1}{50000}}{\frac{1}{20}} = 0.0002$$

Very low probability

How to Apply Bayes' Theorem to Data Mining and Datasets?

<i>Tid</i>	<i>Refund</i>	<i>Marital Status</i>	<i>Taxable Income</i>	<i>Evade</i>
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

- *Target class: Evade*
- *Predictor variables: Refund, Status, Income*
- *What is probability of Evade given the values of Refund, Status, Income?*

$$P(E | R, S, I)$$

Above .5? Predict YES, else predict NO.

How to Apply Bayes' Theorem to Data Mining and Datasets?

<i>Tid</i>	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

- How to compute? $P(E | R, S, I)$
- Need test instance:
 - ▣ What are values of R, S, I ?
 - ▣ Test instance is:
 - Refund=Yes
 - Status=Married
 - Income=60K

Issue: we don't have any training example that these same three attributes values.

Naïve Bayes Classifier

- Why called naïve?
 - ▣ Assumes that attributes (predictor variables) are conditionally independent.
 - No correlation
 - ▣ Big assumption!
- What is conditionally independent?
 - ▣ Variable X is conditionally independent of Y if the following holds:

$$P(X | Y, Z) = P(X | Z)$$

Conditional Independence

Assuming variables X and Y are conditionally independent, can derive:

“given Z , what is the joint probability of X and Y ?”

$$\begin{aligned} P(X, Y | Z) &= \frac{P(X, Y, Z)}{P(Z)} \\ &= \frac{P(X, Y, Z)}{P(Y, Z)} \cdot \frac{P(Y, Z)}{P(Z)} \\ &= P(X | Y, Z) \cdot P(Y | Z) \\ &= P(X | Z) \cdot P(Y | Z) \end{aligned}$$

Naïve Bayes Classifier

Before (simple Bayes' rule):
Single predictor variable X

$$P(Y | X) = \frac{P(X | Y) \cdot P(Y)}{P(X)}$$

Now we have a bunch of predictor
Variables: $X_1, X_2, X_3, \dots, X_n$

$$P(Y | X_1, X_2, \dots, X_n)$$

$$P(Y | X_1, X_2, \dots, X_n) = \frac{P(X_1, X_2, \dots, X_n | Y) \cdot P(Y)}{P(X_1, X_2, \dots, X_n)}$$

$$P(Y | X) = \frac{P(Y) \prod_{i=1}^d P(X_i | Y)}{P(X_1, X_2, \dots, X_n)}$$

Naïve Bayes Classifier

$$P(Y | X) = \frac{P(Y) \prod_{i=1}^d P(X_i | Y)}{P(X_1, X_2, \dots, X_n)}$$

For binary problems: $P(Y | X) > .5$?
Predict YES, else predict NO.

Example: will compute $P(E=Yes \mid Status, Income, Refund)$ and $P(E=No \mid Status, Income, Refund)$

- Find which one is greater (greater likelihood)

Can compute from training data:

Cannot compute / hard to compute:

$$P(X_1, X_2, \dots, X_n)$$

$$P(E = Yes) \quad P(E = No) \quad P(Y)$$

$$P(\text{Refund}=\text{yes}, \text{Status}=\text{married}, \text{Income}=120k)$$

$$P(\text{Refund} = No \mid E = Yes)$$

$$P(X_1 | Y) \quad P(X_3 | Y)$$

- Not a problem, since the two denominators will be the same.
- Need to see which numerator is greater.

Estimating Prior Probabilities for the Class target $P(Y)$

<i>Tid</i>	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

□ $P(\text{Evade}=\text{yes})$

□ $= 3/10$

□ $P(\text{Evade}=\text{no})$

□ $= 7/10$

Estimating Conditional Probabilities for Categorical Attributes $P(X_1 | Y)$

<i>Tid</i>	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

□ $P(\text{Refund}=\text{yes} \mid \text{Evade}=\text{no})$

□ $= 3/7$

□ $P(\text{Status}=\text{married} \mid \text{Evade}=\text{yes})$

□ $= 0/3$

▣ *Yikes!*

▣ *Will handle the 0% probability later*

Estimating Conditional Probabilities for Continuous Attributes

$$P(X_1 | Y)$$

<i>Tid</i>	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

□ For continuous attributes:

1. Discretize into bins

2. Two-way split:

■ $(A \leq v)$ or $(A > v)$

Given a Test Record:

$X = (\text{Refund} = \text{No}, \text{Married}, \text{Income} = 75\text{K})$

Full Example

Tid	Refund	Marital Status	Taxable Income	Evade
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

$$P(\text{NO}) = 7/10$$

$$P(\text{YES}) = 3/10$$

$$P(\text{Refund}=\text{YES} \mid \text{NO}) = 3/7$$

$$P(\text{Refund}=\text{NO} \mid \text{NO}) = 4/7$$

$$P(\text{Refund}=\text{YES} \mid \text{YES}) = 0/3$$

$$P(\text{Refund}=\text{NO} \mid \text{YES}) = 3/3$$

$$P(\text{Status}=\text{SINGLE} \mid \text{NO}) = 2/7$$

$$P(\text{Status}=\text{DIVORCED} \mid \text{NO}) = 1/7$$

$$P(\text{Status}=\text{MARRIED} \mid \text{NO}) = 4/7$$

$$P(\text{Status}=\text{SINGLE} \mid \text{YES}) = 2/3$$

$$P(\text{Status}=\text{DIVORCED} \mid \text{YES}) = 1/3$$

$$P(\text{Status}=\text{MARRIED} \mid \text{YES}) = 0/3$$

For taxable income:

$$P(\text{Income}=\text{above } 101\text{K} \mid \text{NO}) = 3/7$$

$$P(\text{Income}=\text{below } 101\text{K} \mid \text{NO}) = 4/7$$

$$P(\text{Income}=\text{above } 101\text{K} \mid \text{YES}) = 0/3$$

$$P(\text{Income}=\text{below } 101\text{K} \mid \text{YES}) = 3/3$$

Given a Test Record:

$X = (\text{Refund} = \text{No}, \text{Married}, \text{Income} = 75\text{K})$

$$P(\text{NO}) = 7/10$$

$$P(\text{YES}) = 3/10$$

$$P(\text{Refund}=\text{YES} \mid \text{NO}) = 3/7$$

$$P(\text{Refund}=\text{NO} \mid \text{NO}) = 4/7$$

$$P(\text{Refund}=\text{YES} \mid \text{YES}) = 0/3$$

$$P(\text{Refund}=\text{NO} \mid \text{YES}) = 3/3$$

$$P(\text{Status}=\text{SINGLE} \mid \text{NO}) = 2/7$$

$$P(\text{Status}=\text{DIVORCED} \mid \text{NO}) = 1/7$$

$$P(\text{Status}=\text{MARRIED} \mid \text{NO}) = 4/7$$

$$P(\text{Status}=\text{SINGLE} \mid \text{YES}) = 2/3$$

$$P(\text{Status}=\text{DIVORCED} \mid \text{YES}) = 1/3$$

$$P(\text{Status}=\text{MARRIED} \mid \text{YES}) = 0/3$$

For taxable income:

$$P(\text{Income}=\text{above } 101\text{K} \mid \text{NO}) = 3/7$$

$$P(\text{Income}=\text{below } 101\text{K} \mid \text{NO}) = 4/7$$

$$P(\text{Income}=\text{above } 101\text{K} \mid \text{YES}) = 0/3$$

$$P(\text{Income}=\text{below } 101\text{K} \mid \text{YES}) = 3/3$$

$$P(X \mid \text{Class}=\text{No}) = P(\text{Refund}=\text{No} \mid \text{Class}=\text{No})$$

$$\times P(\text{Married} \mid \text{Class}=\text{No})$$

$$\times P(\text{Income}=\text{below } 101\text{K} \mid \text{Class}=\text{No})$$

$$= 4/7 \times 4/7 \times 4/7 = 0.1866$$

$$P(X \mid \text{Class}=\text{Yes}) = P(\text{Refund}=\text{No} \mid \text{Class}=\text{Yes})$$

$$\times P(\text{Married} \mid \text{Class}=\text{Yes})$$

$$\times P(\text{Income}=\text{below } 101\text{K} \mid \text{Class}=\text{Yes})$$

$$= 1 \times 0 \times 1 = 0$$

Since $P(X \mid \text{No})P(\text{No}) > P(X \mid \text{Yes})P(\text{Yes})$

Therefore $P(\text{No} \mid X) > P(\text{Yes} \mid X)$

$\Rightarrow \text{Class} = \text{No}$

Smoothing of Conditional Probabilities

- If one of the conditional probabilities is 0, then the entire product will be 0
- *Idea*: Instead use very small non-zeros values, such as 0.00001

$$\text{Original: } P(x_i | y_j) = \frac{n_c}{n}$$

n = # of training examples that have value y_j

n_c = # of examples from class y_j that take on value x_i

Smoothing of Conditional Probabilities

- *Idea:* Instead use very small non-zeros values, such as 0.00001

$$\text{Original: } P(x_i | y_j) = \frac{n_c}{n}$$

$$\text{Laplace: } P(x_i | y_j) = \frac{n_c + 1}{n + C}$$

n = # of training examples that have value y_j

n_c = # of examples from class y_j that take on value x_i

C = # of classes

w/ Laplace Smoothing:

$$P(\text{NO}) = 7/10$$

$$P(\text{YES}) = 3/10$$

$$P(\text{Refund}=\text{YES} \mid \text{NO}) = 4/9$$

$$P(\text{Refund}=\text{NO} \mid \text{NO}) = 5/9$$

$$P(\text{Refund}=\text{YES} \mid \text{YES}) = 1/5$$

$$P(\text{Refund}=\text{NO} \mid \text{YES}) = 4/5$$

$$P(\text{Status}=\text{SINGLE} \mid \text{NO}) = 3/9$$

$$P(\text{Status}=\text{DIVORCED} \mid \text{NO}) = 2/9$$

$$P(\text{Status}=\text{MARRIED} \mid \text{NO}) = 5/9$$

$$P(\text{Status}=\text{SINGLE} \mid \text{YES}) = 3/5$$

$$P(\text{Status}=\text{DIVORCED} \mid \text{YES}) = 2/5$$

$$P(\text{Status}=\text{MARRIED} \mid \text{YES}) = 1/5$$

For taxable income:

$$P(\text{Income}=\text{above } 101\text{K} \mid \text{NO}) = 4/9$$

$$P(\text{Income}=\text{below } 101\text{K} \mid \text{NO}) = 5/9$$

$$P(\text{Income}=\text{above } 101\text{K} \mid \text{YES}) = 1/5$$

$$P(\text{Income}=\text{below } 101\text{K} \mid \text{YES}) = 4/5$$

Given a Test Record:

$X = (\text{Refund} = \text{No}, \text{Married}, \text{Income} = 75\text{K})$

- I $P(X \mid \text{Class}=\text{No}) = P(\text{Refund}=\text{No} \mid \text{Class}=\text{No})$
 $\times P(\text{Married} \mid \text{Class}=\text{No})$
 $\times P(\text{Income}=\text{below } 101\text{K} \mid \text{Class}=\text{No})$
 $= 5/9 \times 5/9 \times 5/9 = 0.1715$
- I $P(X \mid \text{Class}=\text{Yes}) = P(\text{Refund}=\text{No} \mid \text{Class}=\text{Yes})$
 $\times P(\text{Married} \mid \text{Class}=\text{Yes})$
 $\times P(\text{Income}=\text{below } 101\text{K} \mid \text{Class}=\text{Yes})$
 $= 4/5 \times 1/5 \times 4/5 = 0.128$

Is $P(X \mid \text{No})P(\text{No}) > P(X \mid \text{Yes})P(\text{Yes})$?

$$.1715 \times 7/10 > .128 \times 3/10$$

Therefore $P(\text{No} \mid X) > P(\text{Yes} \mid X)$

$\Rightarrow \text{Class} = \text{No}$

Characteristics of Naïve Bayes Classifiers

- Robust to isolated noise
 - ▣ Noise is averaged out by estimating the conditional probabilities from data
- Handling missing values
 - ▣ Simply ignore them when estimating the probabilities
- Robust to irrelevant attributes
 - ▣ If X_i is an irrelevant attribute, then $P(X_i | Y)$ becomes almost uniformly distributed
 - $P(\text{Refund}=\text{Yes} | \text{YES})=0.5$
 - $P(\text{Refund}=\text{Yes} | \text{NO})=0.5$

Characteristics of Naïve Bayes Classifiers

- Independence assumption may not hold for some attributes
 - ▣ Correlated attributes can degrade performance of naïve Bayes
- But ... naïve Bayes (for such a simple model), still works surprisingly well even when there is some correlation between attributes

References

- *Fundamentals of Machine Learning for Predictive Data Analytics*, 1st Edition, Kelleher et al.
- *Introduction to Data Mining*, 1st edition, Tam et al.
- *Data Mining and Business Analytics with R*, 1st edition, Ledolter